

# students\_performance\_eda\_and\_prediction

December 27, 2025

## 1 Exam Results Analysis and Prediction

This notebook contains an analysis of data concerning students and their exam performance. The goal of the project is to understand the factors influencing exam results and to build a predictive model.

### 1.1 1. Data Loading and Initial Exploration

```
[30]: # Basic imports
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# Load data - change filename to yours
df = pd.read_csv('/content/drive/MyDrive/Students Analysis/
↳Exam_Score_Prediction.csv') # <-- enter the exact filename after upload

# First look
print("First 5 rows:")
print(df.head())

print("\nData Information:")
print(df.info())

print("\nDescriptive Statistics:")
print(df.describe())
```

First 5 rows:

	student_id	age	gender	course	study_hours	class_attendance	\
0	1	17	male	diploma	2.78	92.9	
1	2	23	other	bca	3.37	64.8	
2	3	22	male	b.sc	7.88	76.8	
3	4	20	other	diploma	0.67	48.4	
4	5	20	female	diploma	0.89	71.6	

	internet_access	sleep_hours	sleep_quality	study_method	facility_rating	\
0	yes	7.4	poor	coaching	low	

```

1           yes      4.6      average   online videos      medium
2           yes      8.5      poor       coaching      high
3           yes      5.8      average   online videos      low
4           yes      9.8      poor       coaching      low

  exam_difficulty  exam_score
0            hard      58.9
1      moderate      54.8
2      moderate      90.3
3      moderate      29.7
4      moderate      43.7

```

Data Information:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 20000 entries, 0 to 19999

Data columns (total 13 columns):

#	Column	Non-Null Count	Dtype
0	student_id	20000 non-null	int64
1	age	20000 non-null	int64
2	gender	20000 non-null	object
3	course	20000 non-null	object
4	study_hours	20000 non-null	float64
5	class_attendance	20000 non-null	float64
6	internet_access	20000 non-null	object
7	sleep_hours	20000 non-null	float64
8	sleep_quality	20000 non-null	object
9	study_method	20000 non-null	object
10	facility_rating	20000 non-null	object
11	exam_difficulty	20000 non-null	object
12	exam_score	20000 non-null	float64

dtypes: float64(4), int64(2), object(7)

memory usage: 2.0+ MB

None

Descriptive Statistics:

	student_id	age	study_hours	class_attendance	\
count	20000.000000	20000.000000	20000.000000	20000.000000	
mean	10000.504600	20.473300	4.007604	70.017365	
std	5773.654959	2.284458	2.308313	17.282262	
min	1.000000	17.000000	0.080000	40.600000	
25%	5000.750000	18.000000	2.000000	55.100000	
50%	10000.500000	20.000000	4.040000	69.900000	
75%	15000.250000	22.000000	6.000000	85.000000	
max	20001.000000	24.000000	7.910000	99.400000	

	sleep_hours	exam_score
count	20000.000000	20000.000000

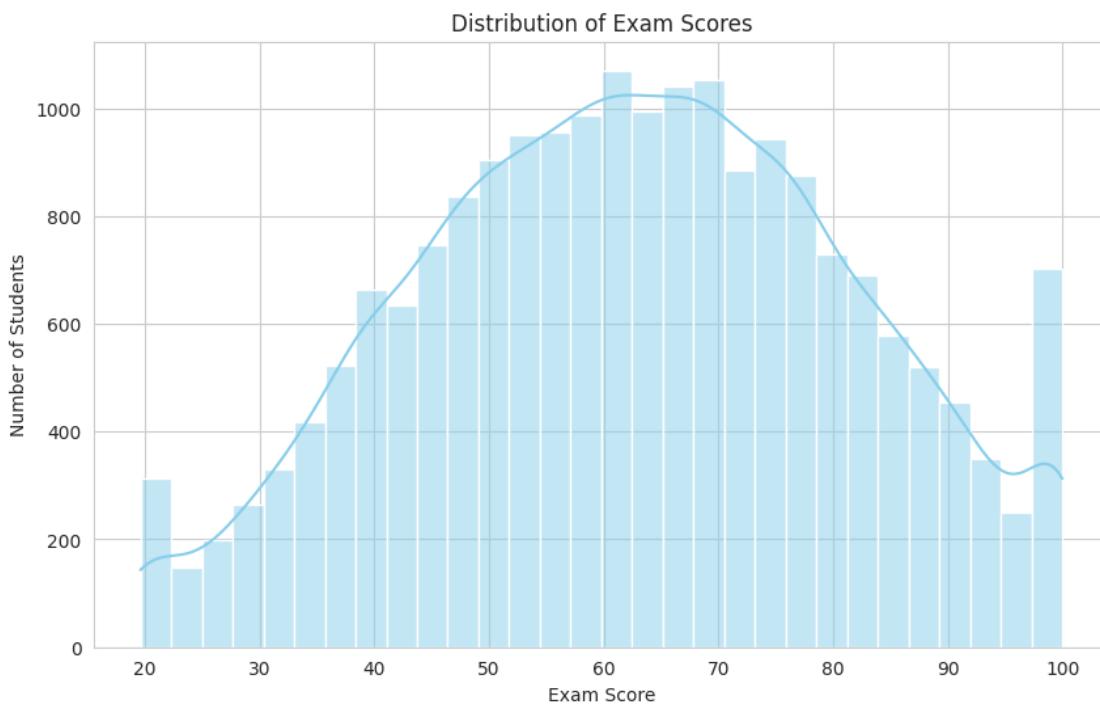
```

mean      7.00856    62.513225
std       1.73209    18.908491
min      4.10000    19.599000
25%      5.50000    48.800000
50%      7.00000    62.600000
75%      8.50000    76.300000
max      9.90000    100.000000

```

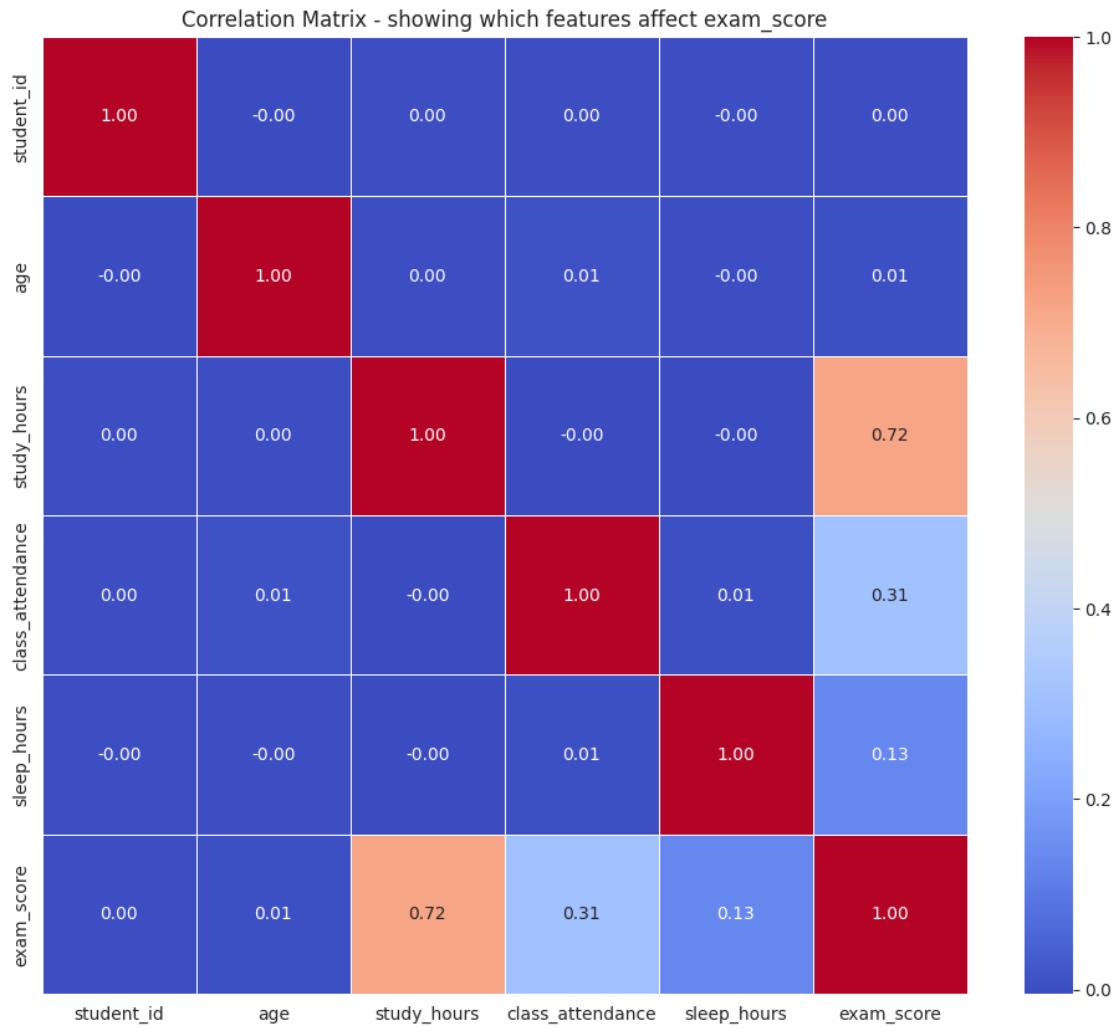
```
[31]: # Set style
sns.set_style("whitegrid")
plt.rcParams['figure.figsize'] = (12, 8)

# 1. Distribution of exam_score (target)
plt.figure(figsize=(10,6))
sns.histplot(df['exam_score'], kde=True, bins=30, color='skyblue')
plt.title('Distribution of Exam Scores')
plt.xlabel('Exam Score')
plt.ylabel('Number of Students')
plt.show()
```



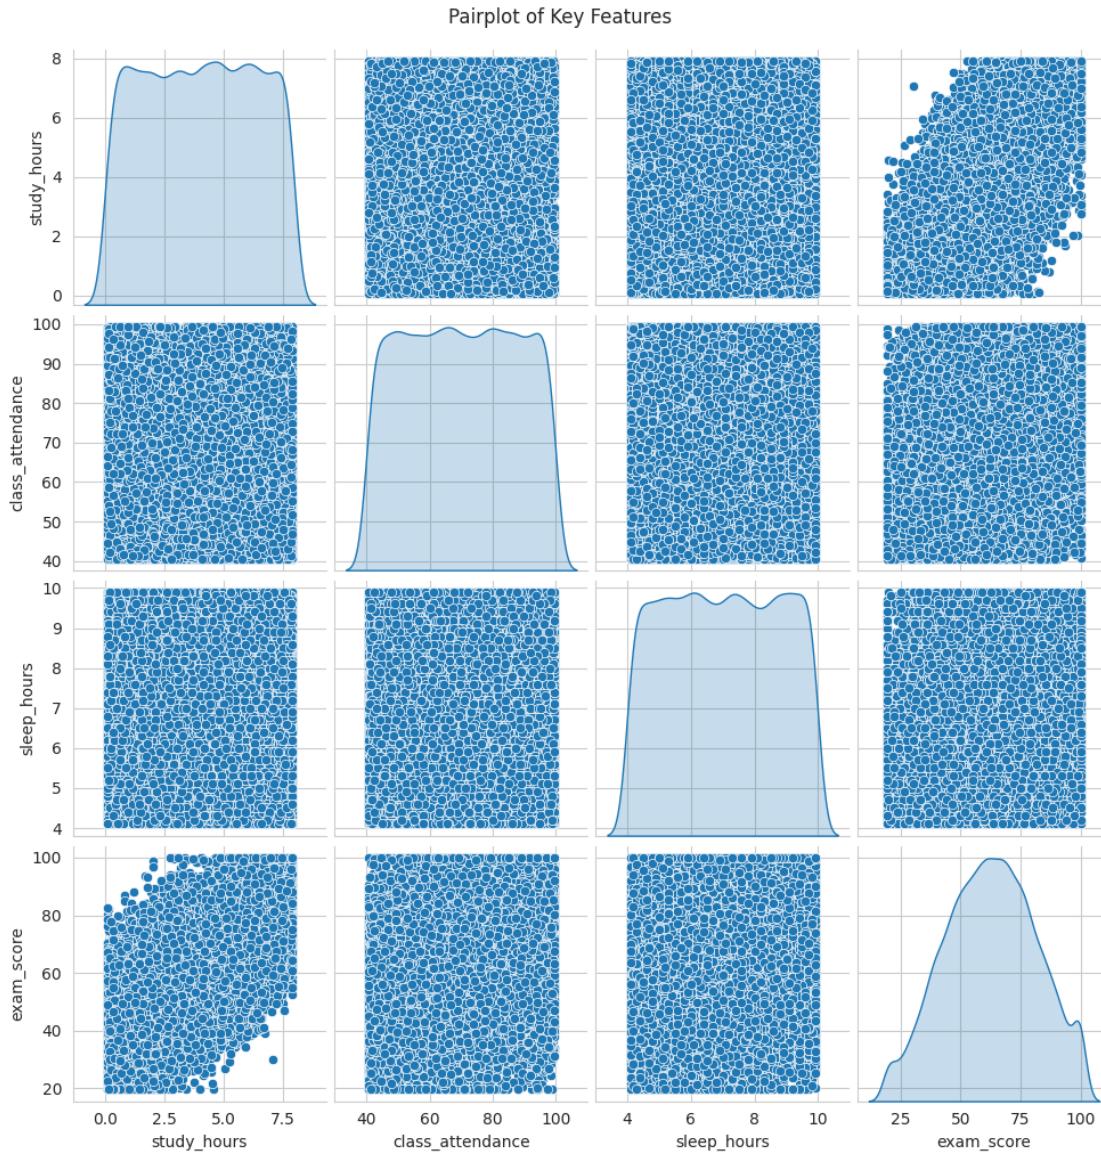
```
[32]: # 2. Correlation heatmap
plt.figure(figsize=(12,10))
corr = df.corr(numeric_only=True)
sns.heatmap(corr, annot=True, cmap='coolwarm', fmt='.2f', linewidths=0.5)
```

```
plt.title('Correlation Matrix - showing which features affect exam_score')
plt.show()
```



[33]: # 3. Pairplot for key features

```
important_cols = ['study_hours', 'class_attendance', 'sleep_hours', ↵
    ↵'exam_score'] # adjust if you have 'previous_scores'
sns.pairplot(df[important_cols], diag_kind='kde')
plt.suptitle('Pairplot of Key Features', y=1.02)
plt.show()
```



## 1.2 2. Data Preparation (Preprocessing)

In this section, we prepare the data for building predictive models. This includes removing unnecessary columns, identifying categorical and numerical columns, as well as splitting the data into training and test sets. We also apply scaling to numerical features and one-hot encoding to categorical columns.

```
[34]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
```

```

# Remove useless columns (student_id, age - zero correlation)
df_clean = df.drop(['student_id', 'age'], axis=1)

# Categorical columns (object) - e.g., gender, course, study_method, etc.
cat_cols = df_clean.select_dtypes(include=['object']).columns.tolist()

# Numerical columns
num_cols = df_clean.select_dtypes(include=['float64', 'int64']).columns.tolist()
num_cols.remove('exam_score') # target

# Target
y = df_clean['exam_score']
X = df_clean.drop('exam_score', axis=1)

# Preprocessing pipeline
preprocessor = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), num_cols),
        ('cat', OneHotEncoder(drop='first', handle_unknown='ignore'), cat_cols)
    ])

# Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

print("Shape of data after preprocessing:")
print(f"X_train: {X_train.shape}, X_test: {X_test.shape}")

```

Shape of data after preprocessing:  
X\_train: (16000, 10), X\_test: (4000, 10)

### 1.3 3. Model Building and Evaluation

```

[35]: from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
import numpy as np

# Full pipeline - preprocessing + model
# 1. Linear Regression as baseline
lr_pipeline = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('model', LinearRegression())
])

# 2. Random Forest - our main model
rf_pipeline = Pipeline(steps=[

```

```

        ('preprocessor', preprocessor),
        ('model', RandomForestRegressor(n_estimators=100, random_state=42))
    )

# Train on training data
print("Training Linear Regression...")
lr_pipeline.fit(X_train, y_train)

print("Training Random Forest...")
rf_pipeline.fit(X_train, y_train)

# Predictions on test data
lr_pred = lr_pipeline.predict(X_test)
rf_pred = rf_pipeline.predict(X_test)

# Metrics
print("\n==== MODEL RESULTS ===")
print("Linear Regression:")
print(f"RMSE: {np.sqrt(mean_squared_error(y_test, lr_pred)):.3f}")
print(f"MAE: {mean_absolute_error(y_test, lr_pred):.3f}")
print(f"R2: {r2_score(y_test, lr_pred):.3f}")

print("\nRandom Forest:")
print(f"RMSE: {np.sqrt(mean_squared_error(y_test, rf_pred)):.3f}")
print(f"MAE: {mean_absolute_error(y_test, rf_pred):.3f}")
print(f"R2: {r2_score(y_test, rf_pred):.3f}")

```

Training Linear Regression...

Training Random Forest...

==== MODEL RESULTS ===

Linear Regression:

RMSE: 9.773

MAE: 7.863

R<sup>2</sup>: 0.733

Random Forest:

RMSE: 10.656

MAE: 8.589

R<sup>2</sup>: 0.683

```
[36]: import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
import numpy as np
```

```

# Only one feature: study_hours
X_simple = df[['study_hours']] # must be a DataFrame, not Series
y = df['exam_score']

# Split (80/20)
X_train_s, X_test_s, y_train_s, y_test_s = train_test_split(X_simple, y, test_size=0.2, random_state=42)

# Model
simple_lr = LinearRegression()
simple_lr.fit(X_train_s, y_train_s)

# Predictions
y_pred_s = simple_lr.predict(X_test_s)

# Metrics
print("== Simple Linear Regression (study_hours only) ===")
print(f"R²: {r2_score(y_test_s, y_pred_s):.4f}")
print(f"RMSE: {np.sqrt(mean_squared_error(y_test_s, y_pred_s)):.3f}")
print(f"MAE: {mean_absolute_error(y_test_s, y_pred_s):.3f}")
print(f"Coefficient (slope): {simple_lr.coef_[0]:.3f} points per study hour")
print(f"Intercept: {simple_lr.intercept_:.3f}")

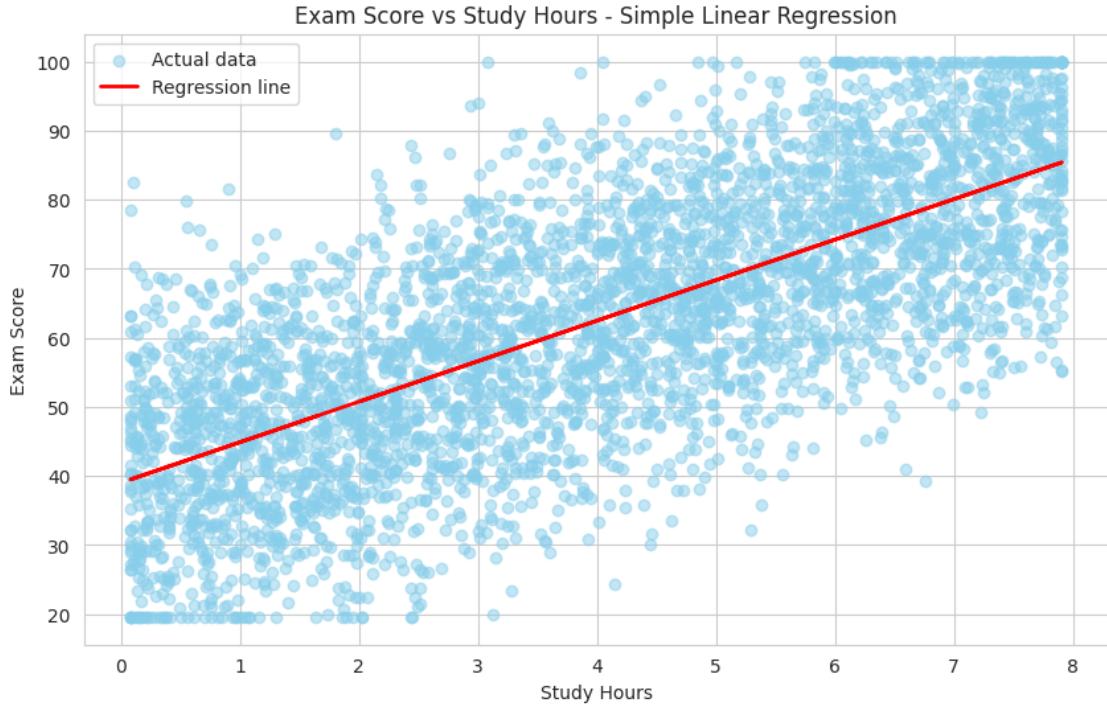
# Visualization
plt.figure(figsize=(10,6))
plt.scatter(X_test_s, y_test_s, alpha=0.5, color='skyblue', label='Actual data')
plt.plot(X_test_s, y_pred_s, color='red', linewidth=2, label='Regression line')
plt.title('Exam Score vs Study Hours - Simple Linear Regression')
plt.xlabel('Study Hours')
plt.ylabel('Exam Score')
plt.legend()
plt.grid(True)
plt.show()

```

```

== Simple Linear Regression (study_hours only) ==
R²: 0.5245
RMSE: 13.042
MAE: 10.641
Coefficient (slope): 5.867 points per study hour
Intercept: 39.020

```



## 1.4 4. Summary and Next Steps

In this section, we present a summary of the results from our analysis and propose potential directions for further research or model improvements.

### 1.4.1 Summary of Results

We conducted an initial analysis of data related to students' exam performance and built predictive models. Key observations:

- **Data Exploration:** Variables such as `study_hours` (hours studied) and `class_attendance` (class attendance) appear to be strongly correlated with `exam_score` (exam score).
- **Predictive Models:**
  - **Simple Linear Regression** (using only `study_hours`) achieved an  $R^2$  of 0.525, RMSE: 13.042, MAE: 10.641. It shows that each additional hour of study translates to approximately 5.87 more points on the exam.
  - **Multivariate Linear Regression** (using multiple features) achieved an  $R^2$  of 0.733, RMSE: 9.773, MAE: 7.863.
  - **Random Forest Regressor** (using multiple features) achieved an  $R^2$  of 0.683, RMSE: 10.656, MAE: 8.589.

It can be seen that the linear regression model with the full set of features (after preprocessing) achieved better results than Random Forest, which is an interesting finding and suggests that for this data, the relationship is linear or close to linear.

### 1.4.2 Next Steps & Possible Improvements

Although the dataset is synthetic (very clean, no outliers, mostly linear relationships), here are realistic ideas for improvement in a real-world scenario:

#### 1. Feature Engineering

Create new features, e.g., interactions (`study_hours × sleep_hours`) or polynomial terms to capture non-linear effects.

#### 2. Advanced Models

Try Gradient Boosting algorithms like XGBoost or LightGBM – they often outperform Random Forest on tabular data.

#### 3. Hyperparameter Tuning

Use `GridSearchCV` or `RandomizedSearchCV` to optimize parameters (e.g., `n_estimators`, `max_depth`).

#### 4. Cross-Validation

Apply k-fold cross-validation for more reliable performance estimates.

#### 5. Real-World Data Handling

In actual projects, focus on dealing with missing values, outliers, and class imbalance – issues absent in this synthetic dataset.

This project demonstrates a complete ML pipeline: EDA → Preprocessing → Modeling → Evaluation. Thanks for swimming through it with me!

[ ]: